

# SIMD-Optimized Numerical Library

## Project Plan

---

### 0) Project Quick Pitch

Goal: A small, clean C++23 or Rust library that delivers measurable 2–10× CPU speedups on core numeric kernels (matmul, dot, conv/fft optional), with clear APIs, correctness tests, and reproducible benchmarks.

Targets: x86-64 (AVX2, AVX-512 if available) and ARM64 (NEON). Falls back to scalar code when SIMD not available.

### 1) Repo Layout

```
simd-lib/
├── include/ or src/      # public headers (C++) / lib modules (Rust)
├── simd/                 # arch-specific intrinsics & dispatch
│   ├── x86/
│   └── arm/
├── kernels/              # matmul, dot, conv, reductions
├── bench/                # micro & macro benchmarks
├── tests/                # unit & property tests
├── data/                 # tiny fixtures for tests
├── cmake.toml/cargo.toml # build config
└── docs/                 # design notes, results, plots
```

### 2) Initial Scope (MVP)

Level-1 kernels (easiest wins)

- dot, axpy ( $y := a \cdot x + y$ ), sum, mean, variance

Level-2 kernels

- gemv (matrix–vector), row/col reductions, sliding-window ops

Level-3 kernels (headline demo)

- gemm (matrix–matrix, blocked + vectorized)

Stretch: 1D conv, FFT (radix-2), or softmax.

### 3) Public API (C++ or Rust)

C++ Example:

```
void dot_f32(const float* x, const float* y, size_t n, float* out);  
void gemm_f32_dispatch(const MatrixView& A, const MatrixView& B, MatrixView C);
```

Rust Example:

```
pub fn dot_f32(x: &[f32], y: &[f32]) -> f32;  
pub fn gemm_f32_dispatch(a: &Matrix, b: &Matrix) -> Matrix;
```

### 4) Architecture + Dispatch

- Compile-time feature sets: SCALAR, SSE2, AVX2, AVX512F, NEON.
- Runtime CPU detection: `cuid/getauxval` → function pointer table.
- Alignment: 32/64-byte alignment; safe tail handling.
- Layout: row-major with stride; blocked GEMM.

### 5) Kernel Design Notes

Dot/Reductions: unroll 4–8 lanes; accumulate in multiple registers; horizontal add at end.

GEMM: blocking, packing, micro-kernel (8x8 AVX2), prefetching, edge handling.

Numerics: FMA where available; optional Kahan summation.

### 6) Portability Strategy

Write scalar reference first, then add intrinsics (`immintrin.h` for x86, `arm_neon.h` for ARM).

### 7) Benchmarking Plan

Micro-benchmarks: dot, axpy, reduce for  $n = 1e3 \dots 1e8$ .

Macro-benchmarks: GEMM for  $256 \times 256 \times 256$ ,  $1024 \times 1024 \times 1024$ , tall/skinny matrices.

Metrics: wall-time, GB/s, GFLOP/s, cache misses (perf/VTune).

Expected speedups: dot/axpy 3–8×, GEMM 5–10×.

## 8) Correctness & Tests

Unit tests vs scalar within tolerance. Randomized property tests. Sanitizer runs.

## 9) Developer Tooling

Profiling: perf, VTune. CI: scalar tests + SIMD gating. Docs with charts & CPU info.

## 10) Milestone Plan (4 Weeks)

Week 1: Scalar dot/axpy; AVX2 dot; benchmark baseline.

Week 2: SIMD axpy/reduce; runtime dispatch; tolerance tests.

Week 3: Scalar blocked GEMM; AVX2 micro-kernel; first benchmarks.

Week 4: Edge handling; docs & charts; optional stretch kernel.

## 11) Stretch Goals

FFT, 1D conv, multithreading, aligned memory pool.

## 12) Show Your Work Checklist

Flamegraphs, benchmark plots, blog-style write-up, reproducible bench scripts.

## 13) Starter Tasks

Set up project; implement scalar dot\_f32; add AVX2 dot; runtime dispatch; benchmark speedup.